

The Stencil Method

Mersimoski Kjanija

[KjMercy/HPC_Assignment](#)



Parallelization of Heat Equation

- Stencil computation for 2D heat diffusion
- MPI --> domain decomposition
- Further optimization with OpenMP

Serial

- 5 Point Stencil
- **Update rule:** weighted combination of self and neighbours
- Heat sources periodically inject energy
- **Naive solution:** double loop over all points N times

MPI

- Domain decomposition
 - Subdomain per task
 - Communication for data exchange
- Communication-computation **overlap**
- Non-blocking calls for neighbouring communication

Overlap

```
inject_energy(...);  
  
//...  
  
fill_buffers(&planes[current], buffers, planes[current].size, neighbours);  
perform_halo_comms(buffers, neighbours, &myCOMM_WORLD, reqs, planes[current].size);  
  
update_inner_plane(&planes[current], &planes[!current]);  
  
//...  
  
MPI_Waitall(8, reqs, statuses);  
  
copy_halo_data(&planes[current], buffers, planes[current].size, neighbours);  
update_border_plane(periodic, N, &planes[current], &planes[!current]);
```

Non Blocking Communications

```
void perform_halo_comms(buffers_t *buffers, int *neighbours, MPI_Comm *comm, MPI_Request *reqs, vec2_t size)
{
    for (uint i = 0; i < 8; i++)
    {
        reqs[i] = MPI_REQUEST_NULL;
    }

    if (neighbours[NORTH] != MPI_PROC_NULL)
    {
        MPI_Irecv(buffers[RECV][NORTH], size[_x_], MPI_DOUBLE, neighbours[NORTH], 0, *comm, &reqs[0]);
        MPI_Isend(buffers[SEND][NORTH], size[_x_], MPI_DOUBLE, neighbours[NORTH], 1, *comm, &reqs[1]);
    }

    if (neighbours[SOUTH] != MPI_PROC_NULL)
    {
        MPI_Irecv(...);
        MPI_Isend(...);
    }

    if (neighbours[EAST] != MPI_PROC_NULL)
    {
        //...
    }
    if (neighbours[WEST] != MPI_PROC_NULL)
    {
        //...
    }
}
```

OpenMP

```
#pragma omp parallel for collapse(2) schedule(static)
  for (uint j = 2; j <= ysize - 1; j++)
    for (uint i = 2; i <= xsize - 1; i++)
      new[IDX(i, j)] = stencil_computation(old, fysize, i, j);
```

```
#pragma omp parallel for schedule(static)
  for (uint j = 1; j <= ysize; j++)
  {
    new[IDX(1, j)] = stencil_computation(old, fysize, 1, j);           // left border
    new[IDX(xsize, j)] = stencil_computation(old, fysize, xsize, j); // right border
  }

#pragma omp parallel for schedule(static)
  for (uint i = 1; i <= xsize; i++)
  {
    new[IDX(i, 1)] = stencil_computation(old, fysize, i, 1);         // top border
    new[IDX(i, ysize)] = stencil_computation(old, fysize, i, ysize); // bottom border
  }
```

OpenMP

```
#pragma omp parallel for reduction(+ : totenergy) collapse(2) schedule(static)
    for (int j = 1; j <= ysize; j++)
        for (int i = 1; i <= xsize; i++)
            totenergy += data[IDX(i, j)];
```

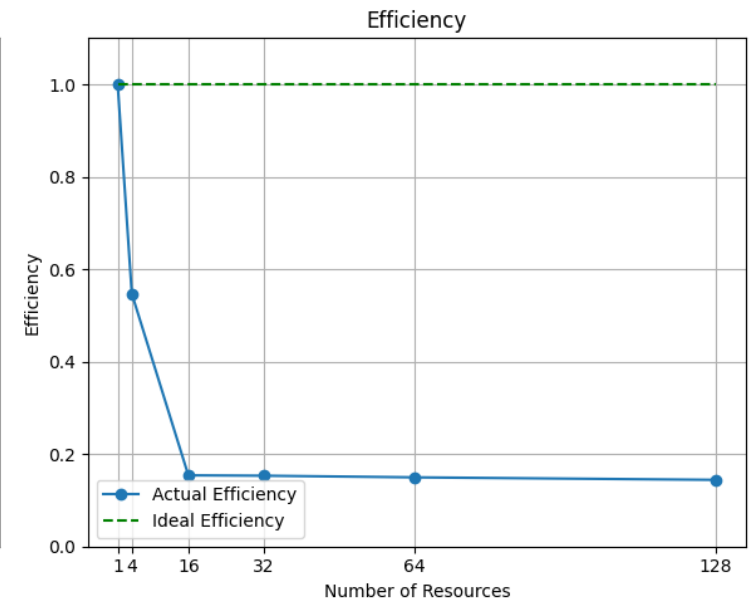
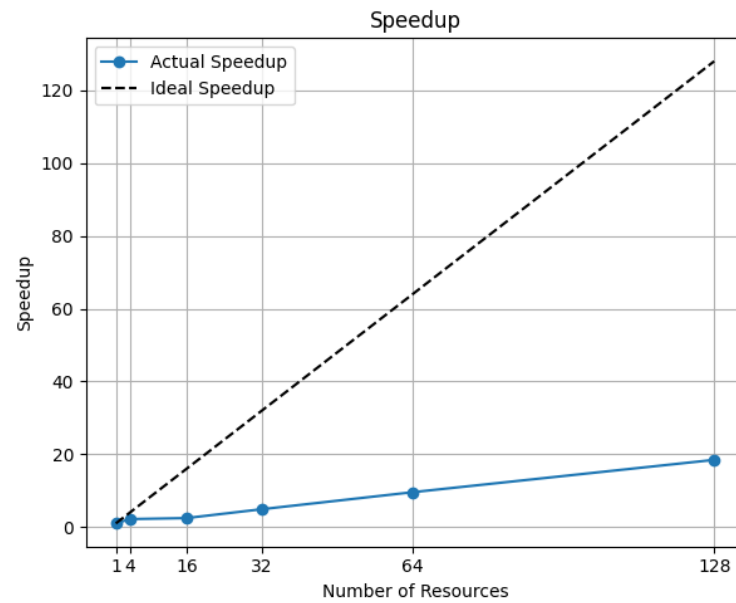
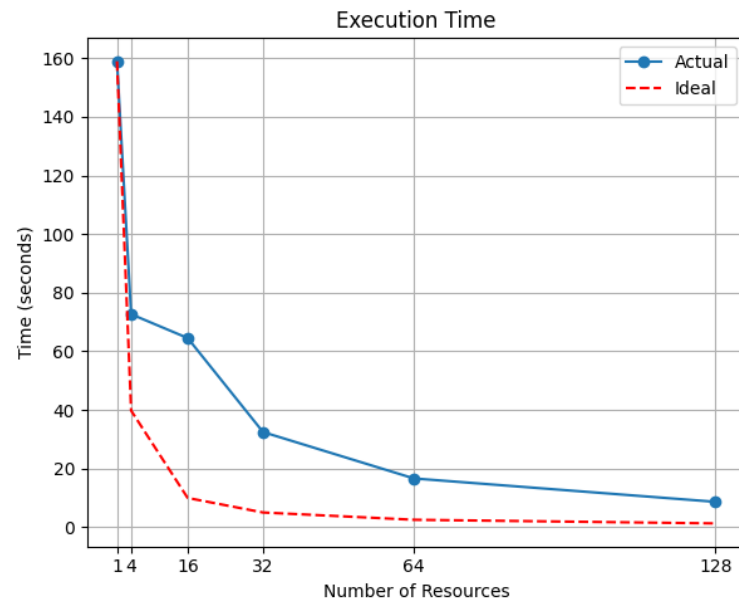

Vectorization

```
#pragma omp parallel for
for (uint j = 2; j <= ysize - 1; j++)
{
    uint i = 2;
    for (; i <= simd_end; i += simd_width) // SIMD loop
    {
        __m256d center = _mm256_loadu_pd(&old[IDX(i, j)]);
        __m256d left = _mm256_loadu_pd(&old[IDX(i - 1, j)]);
        __m256d right = _mm256_loadu_pd(&old[IDX(i + 1, j)]);
        __m256d up = _mm256_loadu_pd(&old[IDX(i, j - 1)]);
        __m256d down = _mm256_loadu_pd(&old[IDX(i, j + 1)]);

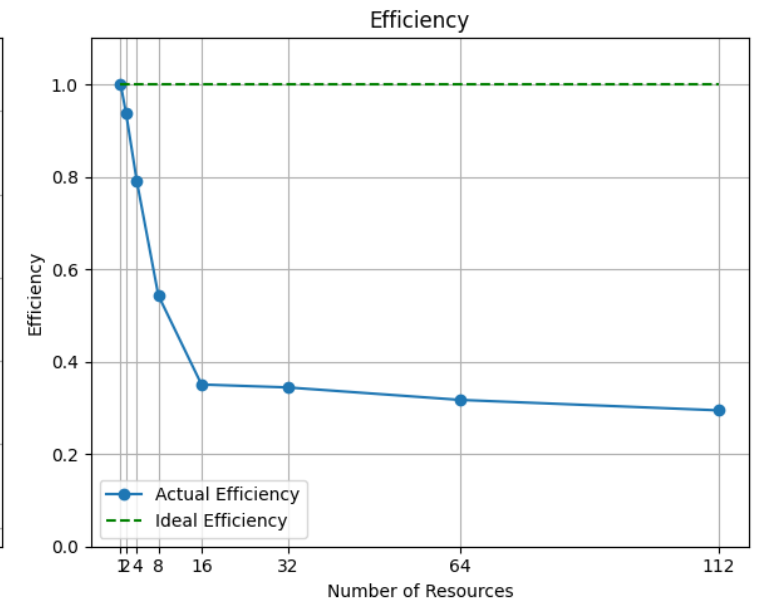
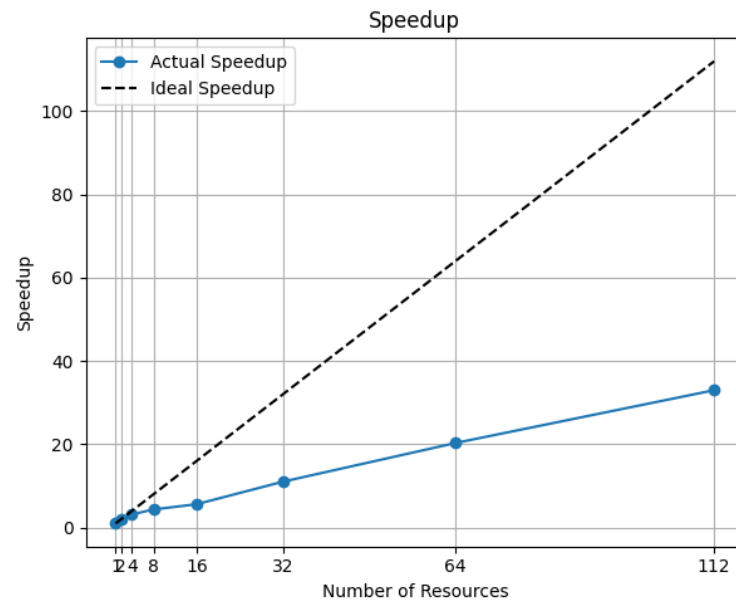
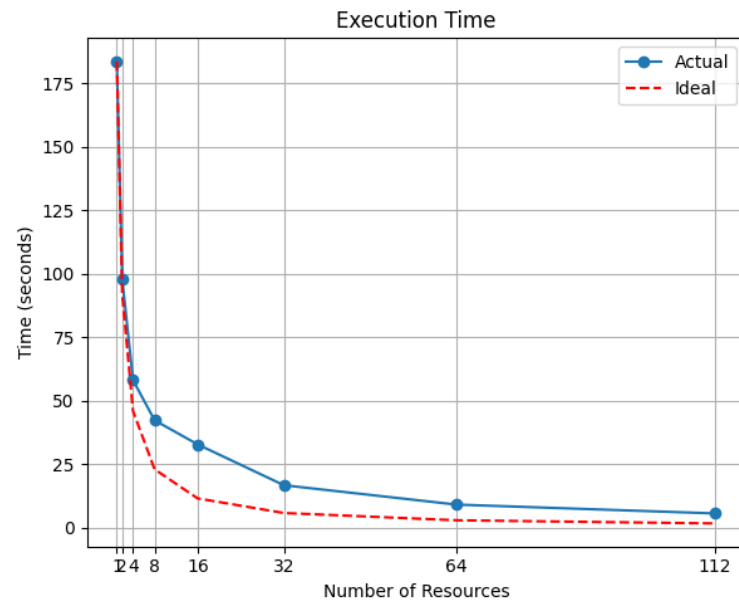
        __m256d res = _mm256_add_pd(_mm256_mul_pd(center, _mm256_set1_pd(0.5)),
                                     _mm256_mul_pd(_mm256_add_pd(_mm256_add_pd(left, right),
                                                                    _mm256_add_pd(up, down)),
                                                                    _mm256_set1_pd(0.125)));
        _mm256_storeu_pd(&new[IDX(i, j)], res);
    }

    for (; i <= xsize - 1; i++)
        new[IDX(i, j)] = stencil_computation(old, fxsize, i, j);
}
```

Threads Scaling - Orfeo



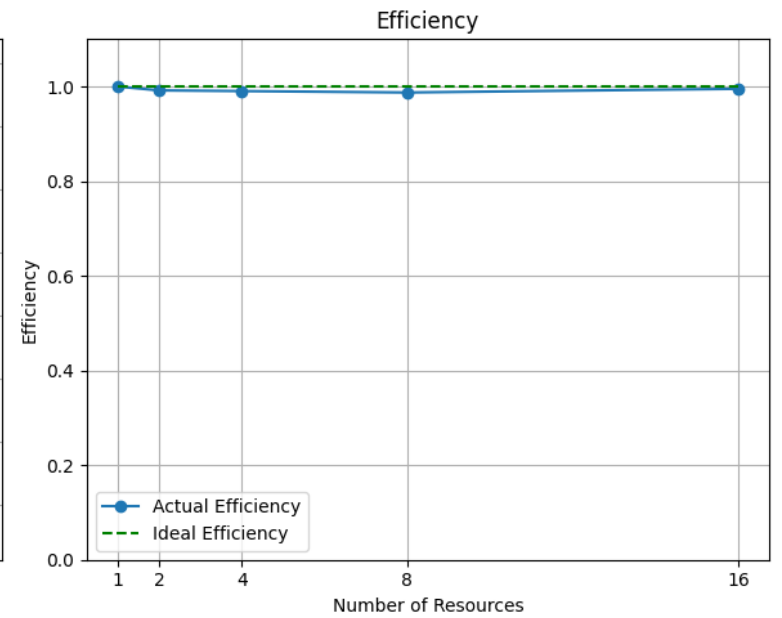
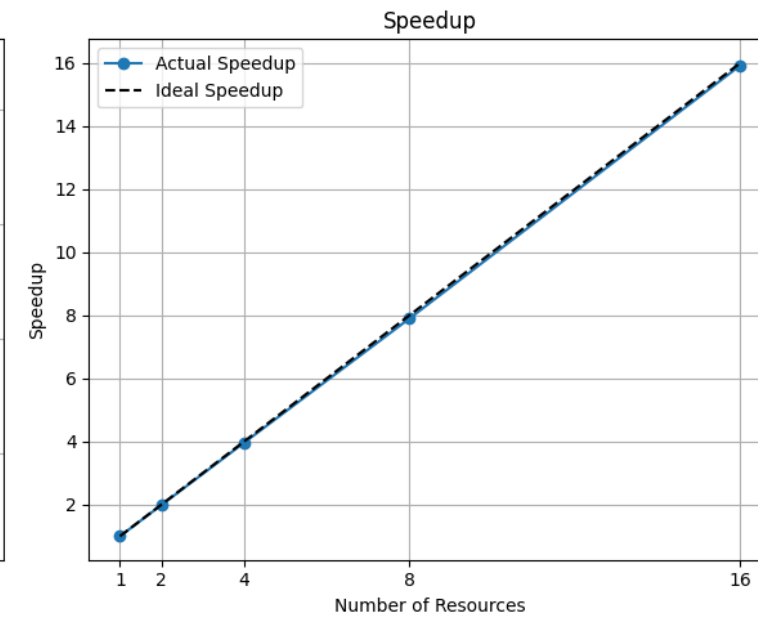
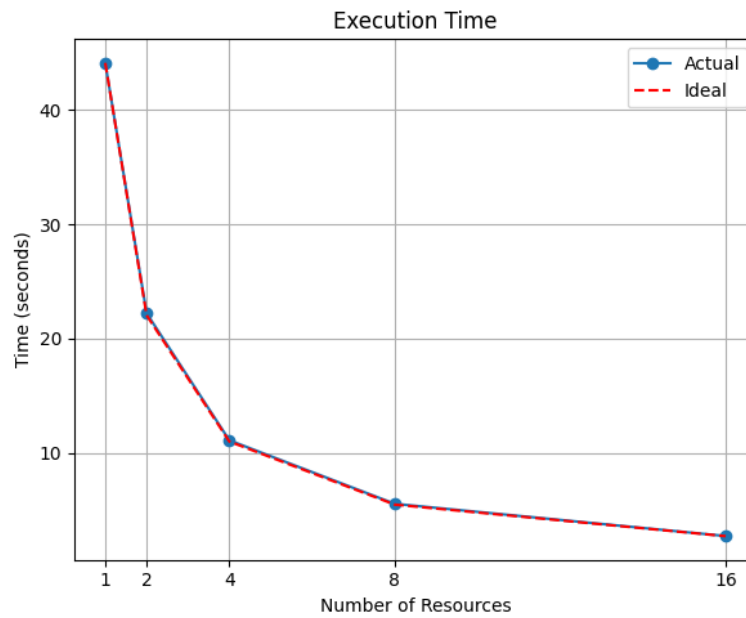
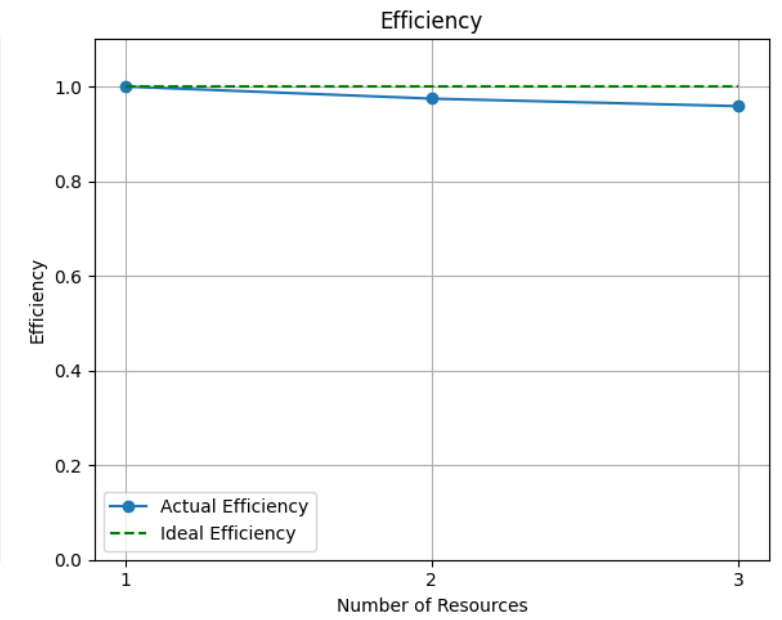
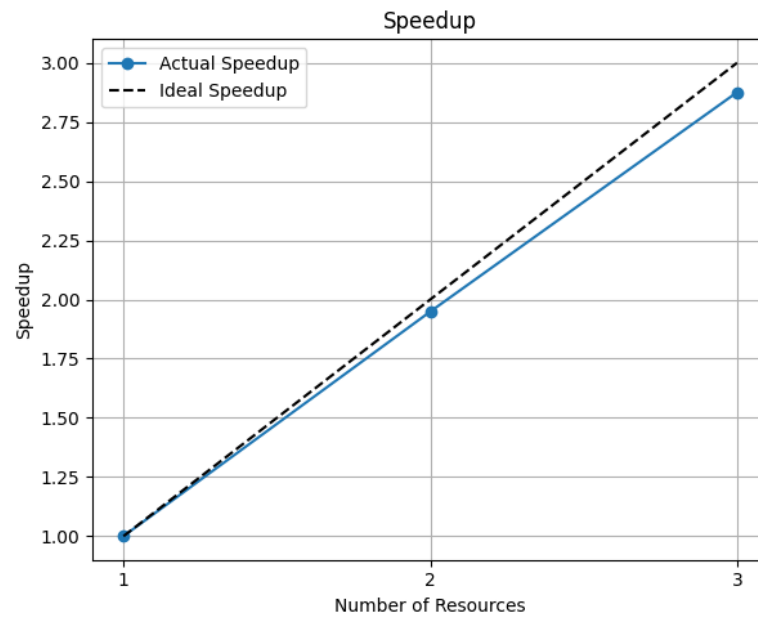
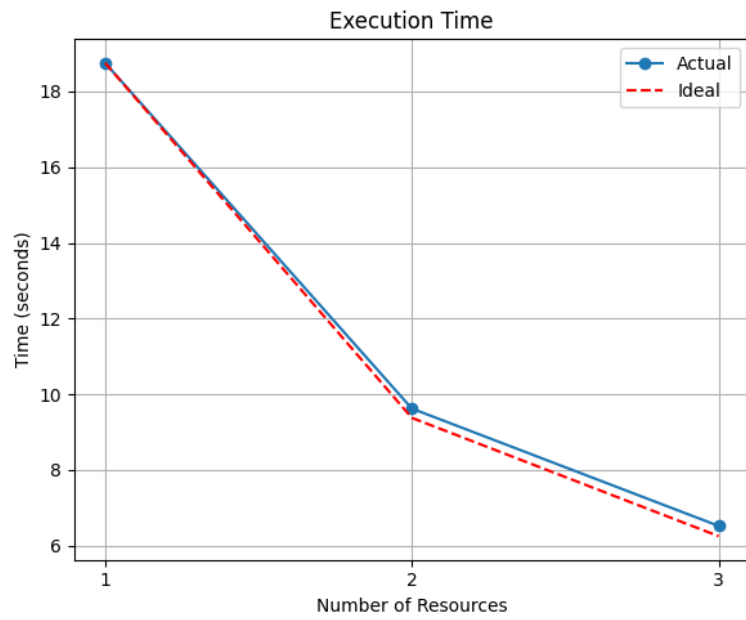
Threads Scaling - Leonardo



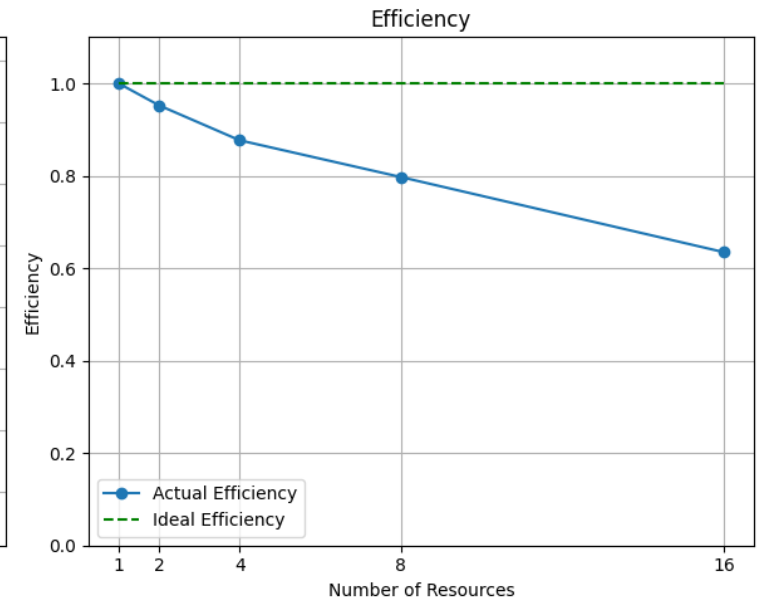
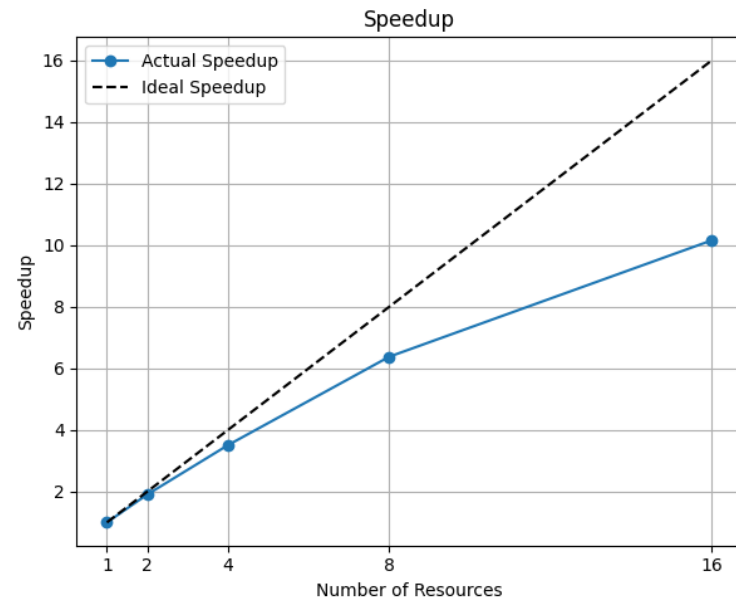
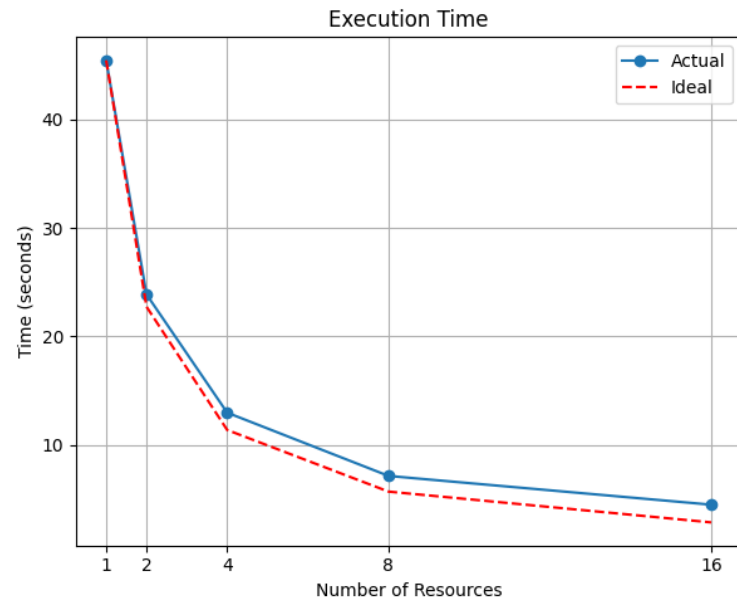
Strong Scaling

- Nodes: 1, 2, 4, 8, 16
- Tasks per node: 8
- CPUs per task: 14

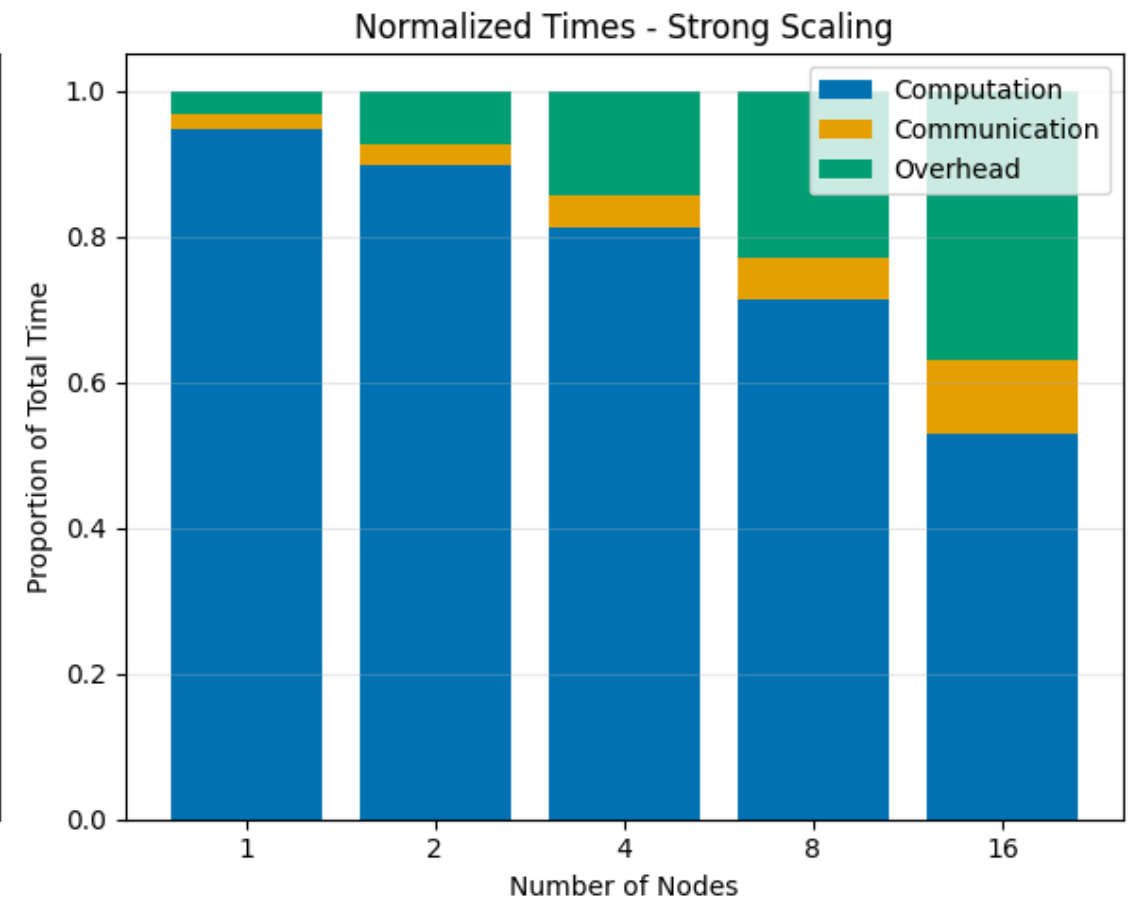
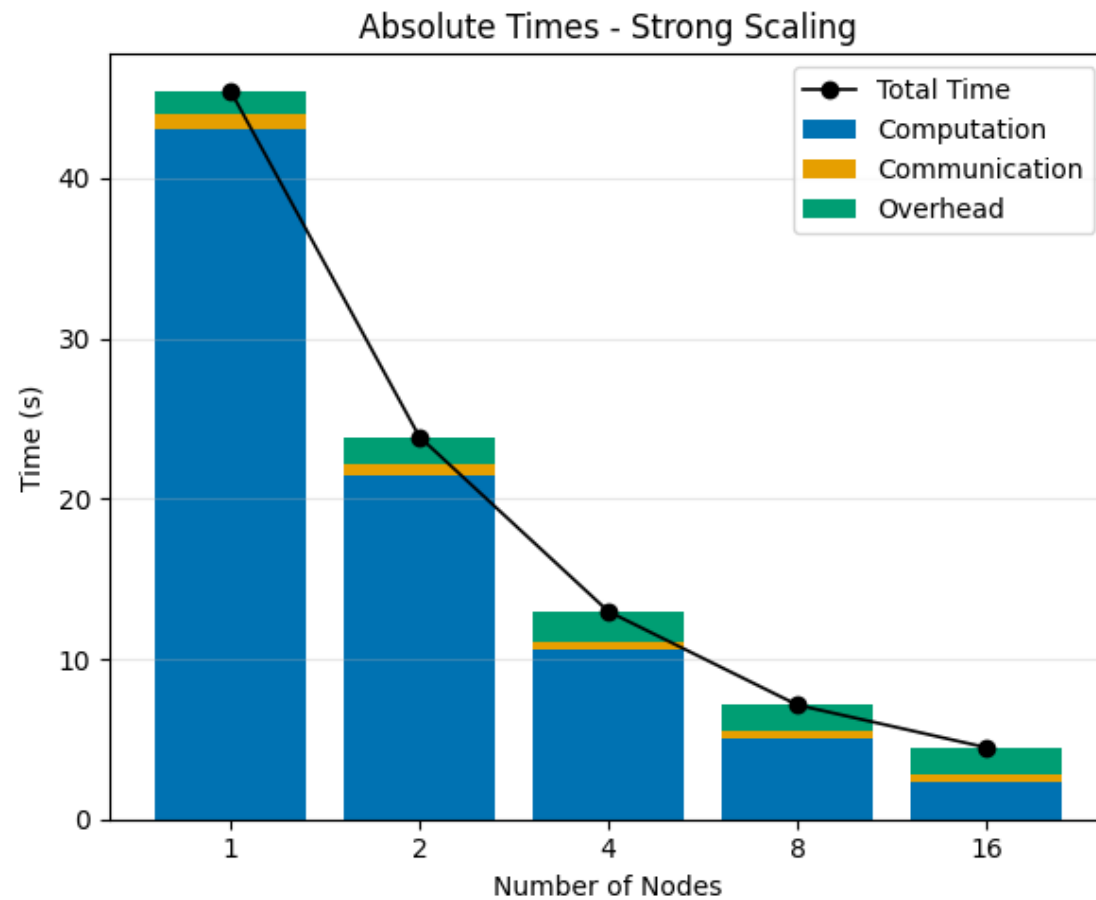
DGCP: 2x56 cores per node.



Strong Scaling - Total time



Time spent doing...



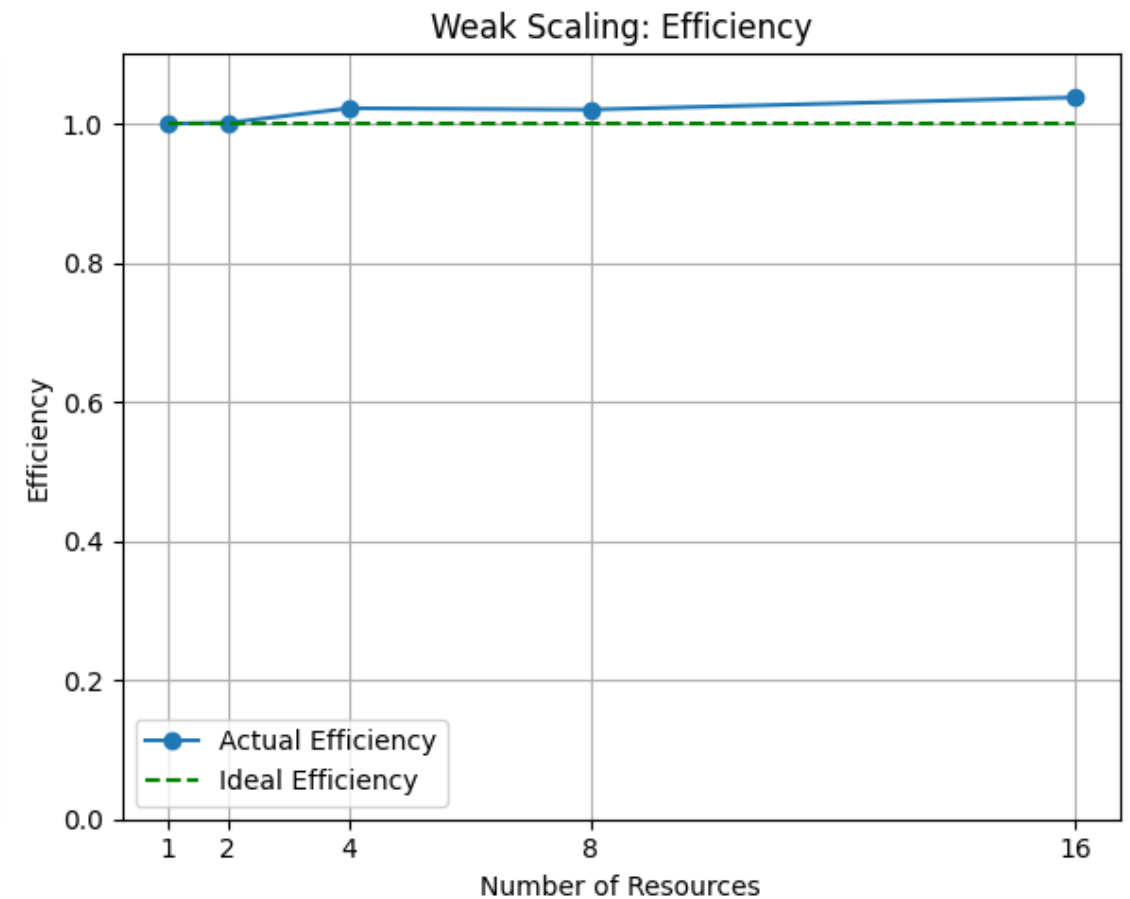
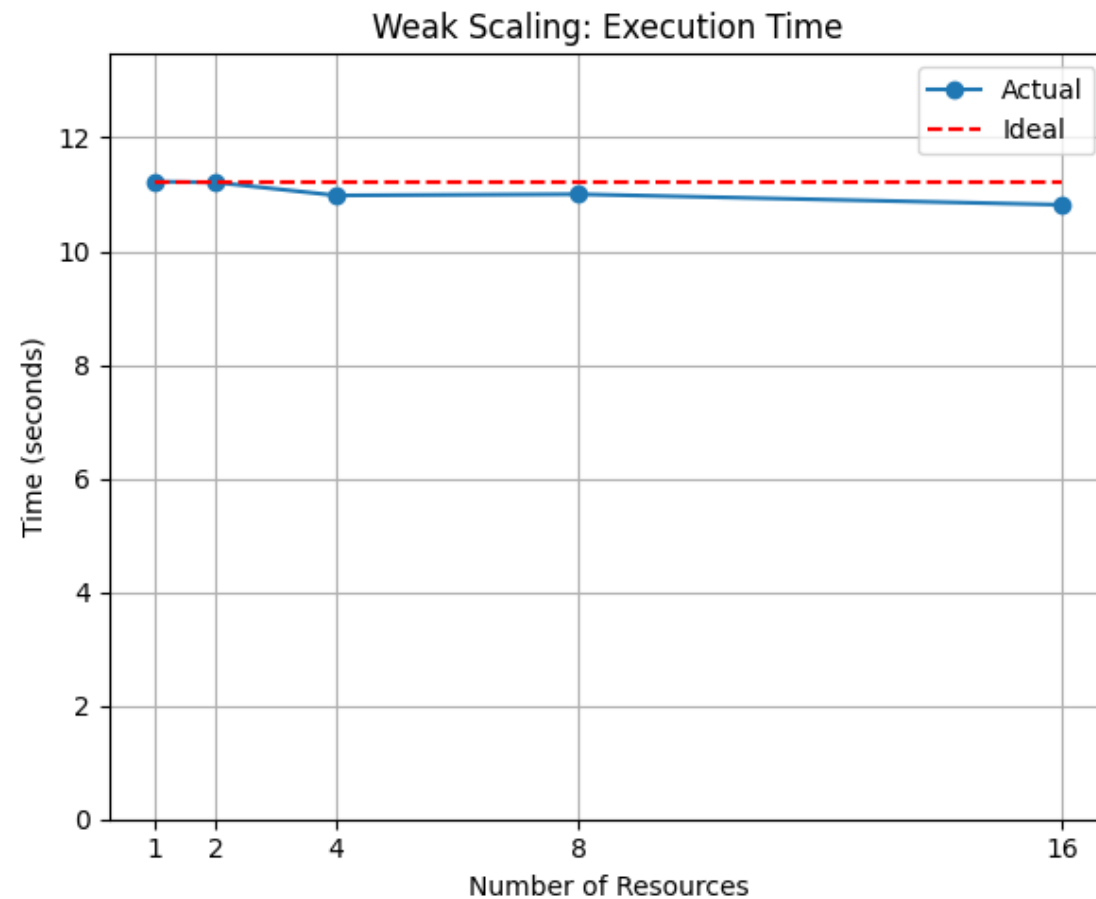
Weak Scaling

- Nodes: 1, 2, 4, 8, 16
- Tasks per node: 8
- CPUs per task: 14

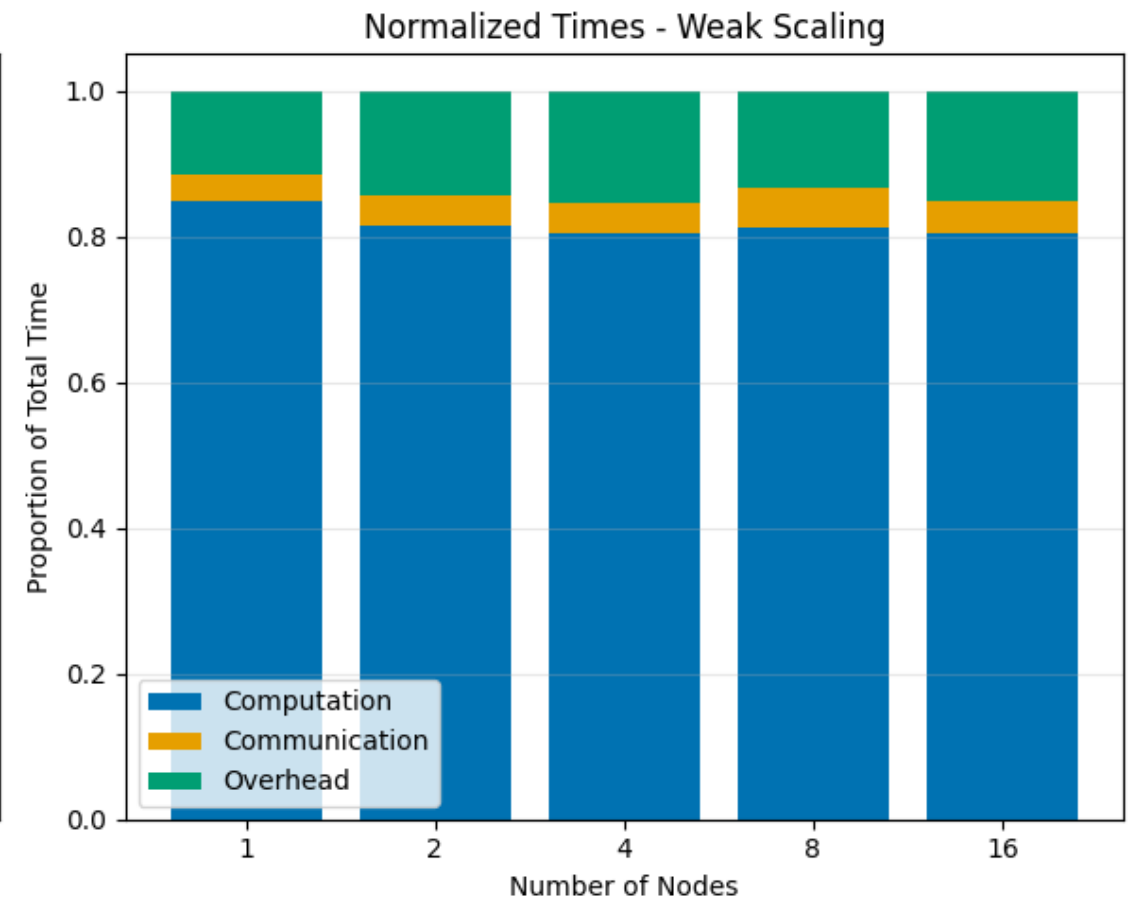
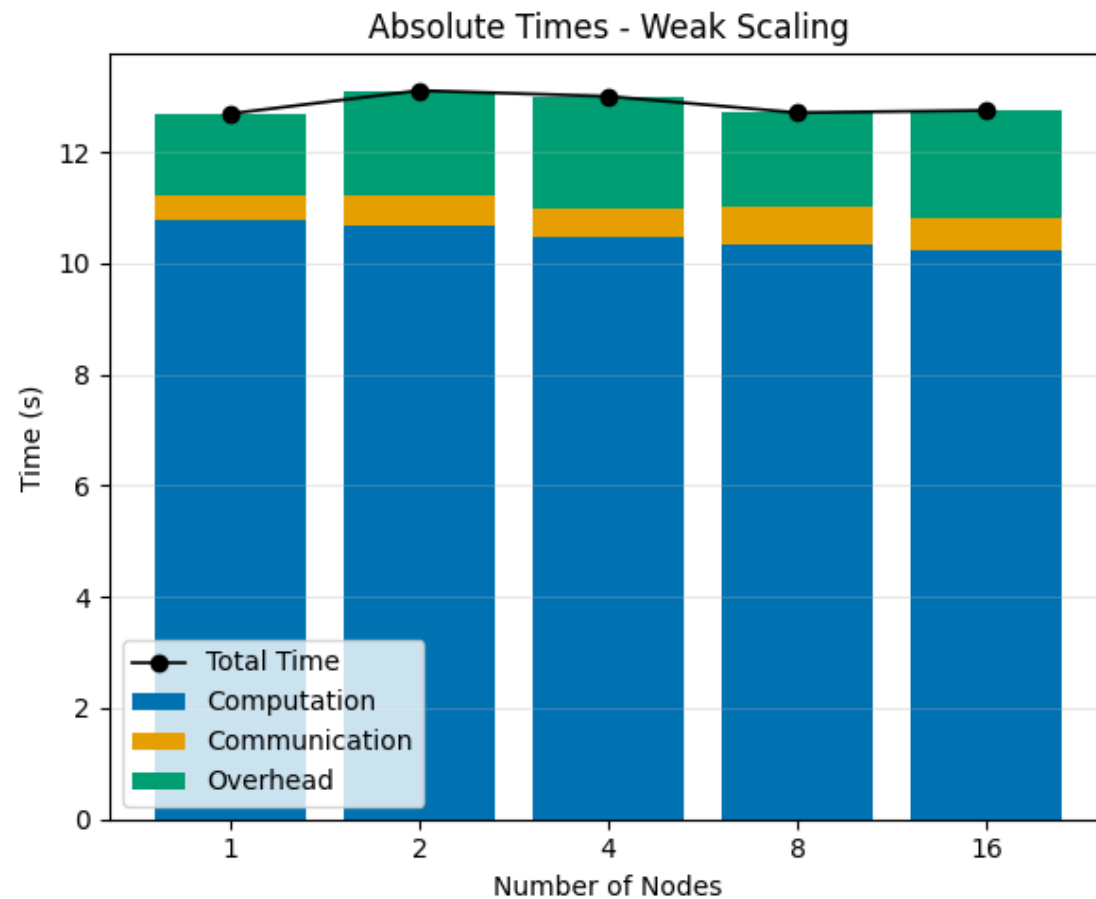
$$x = y = x_0 \times \sqrt{\frac{nt}{nt_0}}$$

Where $nt_0 = 8$ and $x_0 = y_0 = 15000$

Weak Scaling - Total time



Time spent doing...



Valgrind Analysis

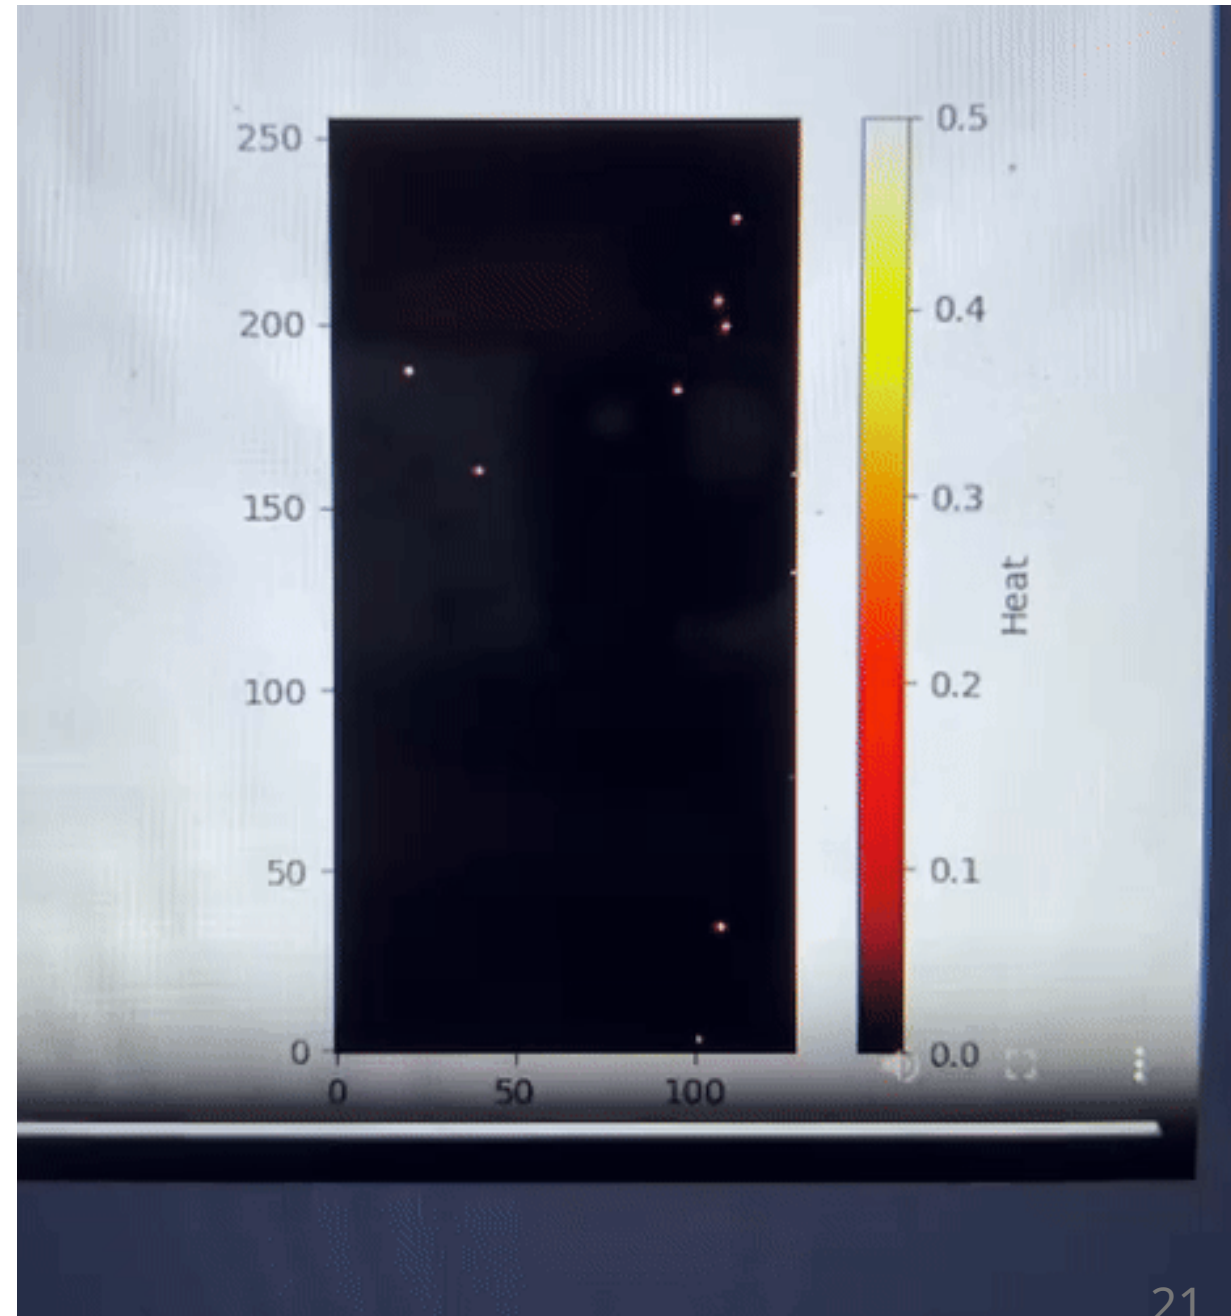
Cache Miss Rates (Cachegrind)

Cache Level	Instructions	Data (Reads)	Data (Writes)	Data (Total)
L1 (L2 run)	0.03%	2.9%	7.8%	4.0%
L1 (L3 run)	0.02%	2.9%	7.8%	4.0%
L2 (LL in run 1)	0.01%	2.3%	7.6%	3.5%
L3 (LL in run 2)	0.00%	0.1%	0.5%	0.2%

Testing and Troubleshooting

- man
- Davide's plot
- GDB
- Assembly

Davide's plot



GDB

```
mpicc -o maindebug -Iinclude src/stencil_template_parallel.c -g  
mpiexec -np 2 gnome-terminal --wait -- gdb -x ./gdb_commands ./maindebug
```

Initial frequent problem: *segmentation fault*.

Solved by: backtracing (`bt` command inside of `gdb`)

Assembly

```
mpicc -O3 -march=native -fopenmp -S -fverbose-asm -masm=intel -Iinclude src/stencil_template_parallel.c -o stencil.s
```

```
    .text
    .globl stencil_parallel
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    vmovsd xmm0, QWORD PTR [rcx+r13*8]    # *_87, *_87
    mov    rbp, r12                      #
    vaddsd xmm0, xmm0, QWORD PTR [rcx+r12*8]    # tmp176, *_87, *_82
# src/stencil_template_parallel.c:315: old[idx - fxsize] + old[idx + fxsize]) *
    mov    r12d, esi                    # i, i
# src/stencil_template_parallel.c:420: new[IDX(i, 1)] = stencil_computation(old, fxsize, i, 1); // top border
    lea    edi, [rdx+1]                #
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    vaddsd xmm0, xmm0, QWORD PTR [rcx+r12*8]    # tmp179, tmp176, *_93
# src/stencil_template_parallel.c:315: old[idx - fxsize] + old[idx + fxsize]) *
    lea    r12d, [rbx+rbp]             # tmp182,
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    lea    r13d, [rbp+0+r8]            # tmp191,
# src/stencil_template_parallel.c:315: old[idx - fxsize] + old[idx + fxsize]) *
    vaddsd xmm0, xmm0, QWORD PTR [rcx+r12*8]    # tmp183, tmp179, *_99
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    lea    r12d, [rsi+r9]              # tmp193,
# src/stencil_template_parallel.c:420: new[IDX(i, 1)] = stencil_computation(old, fxsize, i, 1); // top border
    mov    rdx, rdi                    #
# src/stencil_template_parallel.c:315: old[idx - fxsize] + old[idx + fxsize]) *
    vmulsd xmm0, xmm0, xmm2            # tmp184, tmp183, tmp206
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    vfmadd231sd xmm0, xmm1, QWORD PTR [rcx+rdi*8]    # _103, tmp207, *_76
# src/stencil_template_parallel.c:420: new[IDX(i, 1)] = stencil_computation(old, fxsize, i, 1); // top border
    vmovsd QWORD PTR [r10+rdi*8], xmm0    # *_29, _103
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    vmovsd xmm0, QWORD PTR [rcx+r13*8]    # *_50, *_50
# src/stencil_template_parallel.c:421: new[IDX(i, ysize)] = stencil_computation(old, fxsize, i, ysize); // bottom border
    lea    edi, [r11+rbp]              # tmp189,
# src/stencil_template_parallel.c:314: return old[idx] * 0.5 + (old[idx - 1] + old[idx + 1] +
    vaddsd xmm0, xmm0, QWORD PTR [rcx+r12*8]    # tmp194, *_50, *_55
```

Thank you!